

Machine Learning

CSE204


Jesse Read



Reinforcement Learning

Introduction: Reinforcement Learning

- No supervisor or labels, only **reward** signal
- Reward signal is delayed, sparse, and/or weak
- No i.i.d. 'dataset', but rather an **environment**
- Learning algorithm as part of an **agent**, mapping **observations** to **actions**




Applications

- Games (GO, Backgammon, Atari games, StarCraft, ...)
- Auto-piloting vehicles, robots, ...
- Manage supply and demand (products, electricity)
- Traffic control
- Trading/manage investment portfolio
- Auto-tune the parameters of a deep neural network
- Farming and Agriculture
- Bidding / advertising
- And many others (Health, Sports, Politics, ...)



The number of real-world applications of reinforcement learning is increasing rapidly.

The Agent and the Environment




- 1 Observe the [state of] the environment
- 2 Perform some action
- 3 Obtain reward; and repeat

Goal: take actions to maximize future rewards.




Formulating a Reinforcement Learning Problem

Main considerations:

- What is the state space / what is observed
- What actions are available
- What is the reward function/signal



The agent design (observations, actions, rewards) involves design choices under the constraints of the environment.



The Reward Signal


We observe state s_t and obtain reward r_t , and take action a_t at time t . Over an **episode**, we see

$$s_1, a_1, r_1, \quad s_2, a_2, r_2, \quad \dots \quad , s_T, a_T, r_T$$

Problems:

- Reward may be delayed
- Reward signal may be sparse, and/or weak

i.e., **temporal credit assignment problem**



Fully Observed Environments

If **fully observed**, we see the whole **state**¹. Thus

- \mathcal{S} state space (e.g., $\mathcal{S} = \{A, B, C, D, E\}$)
- \mathcal{A} action space (e.g., $\mathcal{A} = \{1, 2\}$)
- \mathcal{R} reward space (e.g., $\mathcal{R} = \mathbb{R}$)


¹In the partially observed case, we have only an **observation** of a state

Fully Observed Environments

If **fully observed**, we see the whole **state**¹. Thus

- \mathcal{S} state space (e.g., $\mathcal{S} = \{A, B, C, D, E\}$)
- \mathcal{A} action space (e.g., $\mathcal{A} = \{1, 2\}$)
- \mathcal{R} reward space (e.g., $\mathcal{R} = \mathbb{R}$)

For example, Vacuum Cleaner World; There are 8 possible states, 3 possible actions (L,R,S), reward 1 iff cleaned (at end state):



Question: Describe a possible **episode** in this environment.

¹In the partially observed case, we have only an **observation** of a state


Deterministic Environments

In deterministic environments, given an action the next state is known for certain.

For example,

- Noughts & crosses (tic tac toe)
- Chess
- 8-puzzle

×		○
	○	×
×		○



1	3	5
7	2	4
8		6

For example, the 8-puzzle,

1	3	5
7	2	4
8		6

 \Rightarrow

1	2	3
4	5	6
7	8	

It's just a **search** (each **node** is a **state**):

- 1 Generate the search tree to goal state
- 2 Apply payoff to leaf
- 3 Backup
- 4 Choose the best branch

For example, the 8-puzzle,

1	3	5
7	2	4
8		6

 \Rightarrow

1	2	3
4	5	6
7	8	

It's just a **search** (each **node** is a **state**):

- 1 Generate the search tree to goal state
- 2 Apply payoff to leaf
- 3 Backup
- 4 Choose the best branch

But prohibitive if search space is huge/long route to goal (reward)!

Stochastic Environments

In stochastic environments, the agent takes action a from state s and the next state is $s_{t+1} \sim P(\cdot | s_t = s, a_t = a)$.

For example, in the Frozen Lake environment, the ice is slippery and the agent does not always move in the intended direction.

S			
	H		H
			H
H			G

Start, Goal and Hole in the ice

Not a search path from start to goal!

A Markov Decision Process / Model of the Environment

Most reinforcement learning problems can be framed as an MDP:

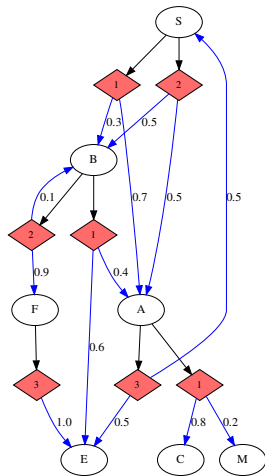
- \mathcal{S} state space (e.g., $\{S, B, F, A, E, C, M\}$)
- \mathcal{A} action space (e.g., $\{1, 2\}$)
- $\mathcal{P}(s'|s, a)$ transition function
- $\mathcal{R}(s, a, s')$ reward function

where \mathcal{P} embodies the Markov property.

We do not necessarily know this model.

How to approach the problem?

We want a policy, of which action is best to take from a given state.



\mathcal{P} shown on edges;
 $\mathcal{R}(A, 1, M) = 1$ and 0 elsewhere

Markov Property

MDPs provide a framework for modeling sequential decision making in stochastic environments.

MDP = Markov Chain + One-step Decision Theory

Markov property: The effects of action a_t from state s_t depend only on those values;

$$P(s_{t+1}|s_t, a_t) = P(s_{t+1}|s_1, \dots, s_t, a_1, \dots, a_t)$$

If $P(s_{t+1}|s_t, a_t) = P(s_{t+1}|s_t)$ we reduce to a Markov chain.

The Policy

A policy

$$\pi : \mathcal{S} \mapsto \mathcal{A}$$

indicates the **action** to take in a given **state**/under the current observation; i.e., the agent observes s_t and takes action $a_t = \pi(s_t)$.

Thus, a policy **defines the behaviour of the agent**.

The Policy

A policy

$$\pi : \mathcal{S} \mapsto \mathcal{A}$$

indicates the **action** to take in a given **state**/under the current observation; i.e., the agent observes s_t and takes action $a_t = \pi(s_t)$.

Thus, a policy **defines the behaviour of the agent**.

Which policy is best?

The Policy

A **policy**

$$\pi : \mathcal{S} \mapsto \mathcal{A}$$

indicates the **action** to take in a given **state**/under the current observation; i.e., the agent observes s_t and takes action $a_t = \pi(s_t)$.

Thus, a policy **defines the behaviour of the agent**.

Which policy is best?

The best policy should take the *best* action from the current state.

$$a_t^* = \pi(s_t)$$

i.e., action a_t to **optimize** ...

The Policy

A policy

$$\pi : \mathcal{S} \mapsto \mathcal{A}$$

indicates the **action** to take in a given **state**/under the current observation; i.e., the agent observes s_t and takes action $a_t = \pi(s_t)$.

Thus, a policy **defines the behaviour of the agent**.

Which policy is best?

The best policy should take the *best* action from the current state.

$$a_t^* = \pi(s_t)$$

i.e., action a_t to **optimize** ... the reward r_t ? No, consider –

- what if the reward is only given at the end (r_T)
- what if a_t has nothing to do with r_{t+1}

We need a payoff function / evaluation metric!

The Return/Gain (Finite Scenario)

The **return** (aka **gain**) (at step t) is thus

$$\begin{aligned} G_t &= \sum_{i=t}^T r_i \\ &= r_t + r_{t+1} + r_{t+2} + \dots + r_T \end{aligned}$$

i.e., the **sum of rewards** until the end of the episode.

The return indicates the **value** of the current state.

It is our 'loss function'; an important consideration ...

- Should it be that $r_{t+1} = r_T$ (as above)?
(or should closer rewards be more important?)
- What is T ? Perhaps $T = \infty$?
(agent may not survive until then)

The Return (Infinite Scenario)

The **return** (at step t) is thus

$$G_t = \sum_{k=0}^{\infty} \gamma^k r_{k+t+1} \quad (1)$$

$$= r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \quad (2)$$

for **discount factor** $\gamma \in (0, 1)$ which indicates the relative value of closer rewards.

The Return (Infinite Scenario)

The **return** (at step t) is thus

$$G_t = \sum_{k=0}^{\infty} \gamma^k r_{k+t+1} \quad (1)$$

$$= r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \dots \quad (2)$$

for **discount factor** $\gamma \in (0, 1)$ which indicates the relative value of closer rewards.

Issue: G_t is from the future (and may include up to r_{∞})!
Instead, consider the **expected return** $\mathbb{E}[G_t]$.

The Value Function

The **value function** (or **state-value function**),

$$V^{\pi}(s) = \mathbb{E}[G_t | s_t = s] \quad (3)$$

maps a state to a **value**.

The value of a state **s** is the **expected return** from that state following policy π .

We may think of a vector of $|\mathcal{S}|$ values;
e.g., if $\mathcal{S} = \{s_1, s_2, s_3, s_4\}$:

	s_1	s_2	s_3	s_4
V				

The **policy** is implicit: **move to states with high value**.

The Action-Value Function

The **action-value function**,

$$Q^{\pi}(s, a) = \mathbb{E}[G_t | s_t = s, a_t = a]$$

maps a state *and action* to a value.

We may think of a table of $|\mathcal{S}| \times |\mathcal{A}|$ values;
e.g., with $\mathcal{A} = \{a_1, a_2\}$:

Q	a_1	a_2
s_1		
s_2		
s_3		
s_4		

The **greedy policy** is explicit:

$$a_t = \pi(s_t) = \operatorname{argmax}_{a \in \mathcal{A}} Q(s_t, a)$$

Given Q (or V) we can construct an optimal policy using dynamic programming; **solved!**

What is the problem? We don't usually have Q or the MDP (observed environment) which can produce it!

That's where reinforcement learning comes in: Estimate Q (hence, e.g., Q-Learning!)

Exploration vs Exploitation

We need to learn the system dynamics through interaction.

One possibility: Use a **behaviour policy** (possibly random) μ :

$$a_t \sim \mu$$

(this is an **off-policy** method).

- 1 Play many episodes with policy μ ; record G_t for each (s_t, a_t)
- 2 Use these samples to approximate the expectation:

$$Q^\mu(s, a) = \frac{1}{n} \sum_{i=1}^n G_t^{(i)} \approx \mathbb{E}[G_t | s, a]$$

- 3 Employ the **greedy policy** π on Q .

Trade-off: Exploration vs exploitation.
--

Another Way: Stochastic Optimization

We can perform a search directly in policy space.

Imagine:

- Neural network as π .
- Try many π_1, \dots, π_M in the environment
- Find the one that performs best

You won't find this covered extensively in reinforcement learning books, because it's already covered in [stochastic optimization](#), including [evolutionary methods](#).

Can be “*embarrassingly effective*” for a small number of parameters (< 100), especially for continuous state spaces, in particular domains such as robotics; but often severely limited on larger problems.

Types of Reinforcement Learning Agents

Value-Based or Policy-Based ?

- Value Based (value function; policy is implicit)
- Policy Based (no value function)
- Actor Critic (both value function + policy)

Model Free vs **Model Based**: We can build a model of the environment, or do without it.




Image from [\[1\]](#)

Summary: Important Concepts

- Agent and Environment
- States, Actions, Reward
- Model of the environment (MDP)
 - Markov property
 - Deterministic vs stochastic
 - Unknown vs fully observable
- Return/Gain (finite horizon vs infinite horizon; γ)
- Policy
- State Value and Action-Value functions
- Exploration vs Exploitation trade-off

Reinforcement learning is difficult (compared to supervised learning), but can be very useful – there are many applications where supervised learning does not apply.

Machine Learning

CSE204

Jesse Read



Reinforcement Learning